

Missing Data - Introductory notes

David Pham

Contents

1	Wikipedia	1
2	Matloff Blog post	2
3	Missing data [@schafer2002missing]	3
4	Data analysis using regression and multilevel/hierarchical models [@gelman2006data]	6
5	Statistical analysis with missing data [@little2002statistical]	7
6	Imputing Missing Data for Gene Expression Array [@hastie1999imputing]	7
7	Missing value estimation methods for DNA microarrays [@troyan-skaya2001missing]	9
8	Datasets	10
9	Chained Equations	10
10	FLAS Completion	10
11	Bibliography	10

These are my current notes and the themes

1 Wikipedia

Case deletion (CC) The method does not introduce any bias if the missing values are uniformly distributed.

Single imputation If one sort the data matrix according some order, *Last observation carried forward* is the method of replacing the missing value with last valid value.

The missing value can also be replaced with the mean of the other observations, however, correlations are attenuated.

Regression imputation use the other variables as predictors to replace the missing value, although precision is misleadingly augmented, hence does not reflect the statistical errors of the missing data. This problem is partially solved by multiple imputation.

Multiple imputation The multiple imputation [Rubin (1987)] is similar to bootstrapping method: Missing variables are simulated, say B times, and the desired statistics are averaged except for the standard error which is constructed by adding the variance of the imputed data and the within variance of each data set.

2 Matloff Blog post

Complete-case analysis (CC), listwise deletion Delete all record for which at least one variable is missing.

Single and multiple imputation Estimation of the distribution of missing variables conditional on the others and then sampling from that distribution. Multiple alternate matrix are generated without the NAs.

In multiple imputation, the distribution of each variable conditional and the others is fitted and in case of missing value, a sample is drawn from this distribution.

Available cases (AC), pairwise deletion Keep the observation if the missing feature is not retained for the desired measure, for example the correlation (where only 2 variables are needed). It can, nonetheless, produce correlations over 1.

2.1 MCAR: Missing Completely at Random

Let Y the variable of interest, $M \in \{0, 1\}$ denotes if Y is missing, and D the other variables than Y . This is often denoted as

$$P(M = 1|Y = s, D = t) = P(M = 1),$$

or equivalently

$$P(Y = s, D = t|M = i) = P(Y = s, D = t), i \in \{0, 1\}.$$

2.2 MAR: Missing at Random

For multiple imputation, one requires only $M \perp Y|D$, that is

$$P(M = 1|Y = s, D = t) = P(M = 1|D = t),$$

2.2.1 Conditional estimation under MAR

In practice, problems arise as D might not hold any predictive ability of the desired variable and that D might as well contain missing data. Interestingly

$$\begin{aligned}
 P(Y = s|D = t, M = i) &= \frac{P(Y = s, D = t, M = i)}{P(D = t, M = i)} \\
 &= \frac{P(Y = s, D = t)P(M = i|Y = s, D = t)}{P(D = t, M = i)} \\
 &= \frac{P(Y = s|D = t)P(D = t)P(M = i|D = t)}{P(D = t, M = i)} \\
 &= P(Y = s|D = t).
 \end{aligned}$$

Hence if we are interested in the relationship between Y and D , that is the conditional distribution Y given D , the fact that it is missing or not will not introduce bias, hence CC and AC would perform equally well. This is ironic as MAR is meant to apply where CC and AC should not be used.

2.2.2 Unconditional estimation under MAR

Observe that

$$P(Y = s|M = 0) = \frac{P(M = 0|Y = s)}{P(M = 0)} P(Y = s),$$

hence our estimation of $P(Y = s)$ might still be biased with the factor of $P(M = 0|Y = s)/P(M = 0)$.

3 Missing data [schafer2002missing]

With or without missing data, the goal of a statistical procedure should be to make valid and efficient inferences about a population of interest—not to estimate, predict, or recover missing observations nor to obtain the same results that we would have seen with complete data.

Let Y_{com} denote the complete data, and denote its partitions with observed and missing data $Y_{com} = (Y_{obs}, Y_{mis})$. If R is the random variable representing missingness, then MAR (also called ignorable non-response) is defined as

$$P(R|Y_{com}) = P(R|Y_{obs}),$$

and MCAR

$$P(R|Y_{com}) = P(R).$$

Missing not at random (MNAR) or non-ignorable non-response, is the situation when MAR is violated. Issue with MAR is, it is often unverifiable, however, only little deviation of estimates and standard errors are observed in practice.

$P(Y_{com}; \theta)$ can be interpreted as either the sampling mechanism of Y_{com} with parameter θ or the likelihood function. The following formula

$$P(Y_{obs}; \theta) = \int P(Y_{com}; \theta) dY_{mis}$$

provides a sampling distribution only when MCAR holds and is a valid likelihood function when MAR is assumed (favoring the Bayesian view). For MNAR, R and an additional parameter ξ defining the distribution of R has to be added:

$$P(Y_{obs}, R; \theta, \xi) = \int P(Y_{com}; \theta) P(R; \xi) dY_{mis}.$$

3.1 Older Methods

Listwise and Pairwise deletion Listwise deletion (case deletion or complete-case analysis) dismiss all observation with any missing values and pairwise deletion (available-case analysis) uses different sets of sample units for different parameters. Critics of AC are that the standard errors or other measures of uncertainty are difficult to assess as the parameters are computed from different sets of units.

CC analysis only works with MCAR but even if it holds, MCAR can be inefficient (e.g. with large data matrix with mild rates of missing values).

Reweighting Reweighting can eliminate bias from CC, for more details [Little2002statistical, chapter 4.4]. It is easy to use with univariate and monotone missing patterns.

Average imputation It replaces the missing value with the mean of the observation. This introduce bias and underestimate the standard errors. The new value is an artifact of a specific data sets and disturbs the scale of the variables. If MI is not feasible, then averaging is a reasonable choice if

reliability is high ($\alpha > 0.7$) and each group of items to be averaged seems to form a single, well, defined domain

3.2 Single imputation

Imputation is the process of predicting the missing value conditional on the other values. It has the advantages of sharing the same data set to all researcher working on a common project. See [Little2002statistical] for shortcomings of single imputation.

Imputing unconditional means Mean substitution consists of replacing the missing value of a variable with the average across all the other non-missing observations. Weakness are that confidence intervals $\bar{y} \pm z_\alpha \sqrt{S^2/N}$ are narrowed by overstating the number of observation N and the downward bias into S^2 . Under MCAR the coverage is only $2\Phi(z_\alpha r) - 1$ where r is the rate of missingness.

Imputing from unconditional distributions Hot deck imputation fills in non-respondents' data with values from actual respondents, that is we replace with a random draw from the observed values. This methods still distort correlation and standard errors.

Imputing conditional means In the univariate situation (where only one value is of interest), one can fill with a prediction from the other variable using regression methods.

This is nearly optimal for a limited class of estimations problem if special correction are made to standard errors.

However, it overstates the correlation and covariance as R^2 for imputed value is 1.00.

Imputing from conditional distribution Under MAR assumption, the weaknesses from the previous methods are overcome by drawing an observation from the fitted regression distribution of Y given X . In general, one has to sample from

$$P(Y_{mis}|Y_{obs}, \theta),$$

where, in practice, we replace θ with its estimated value $\hat{\theta}$ from Y_{obs} . With monotone patterns, one can set a sequence of regression for Y_j given Y_1, \dots, Y_{j-1} , for $j \in 1, \dots, p$.

Undercoverage and reasonable application In a simulation exercise, one can deduce that the actual coverage is much lower than 95%. Compared to CC, if the missing rate is low, single imputation might still be a valid method. For example, if $p = 25$ and the missing rate $r = 0.03$, then CC would delete $1 - (1 - r)^p = 0.53$ of the cases, whereas conditional distribution would allow to use all the participants.

3.2.1 Maximum likelihood estimation

One of the advantage of using the MLE $\hat{\theta}$ is hypothesis testing. If $\tilde{\theta}$ is the MLE for the null hypothesis, one could use likelihood-ratio tests and thus compare

$$2[l(\hat{\theta}; Y_{obs}) - l(\tilde{\theta}; Y_{obs})],$$

and the $(1 - \alpha)$ -quantile of the χ_p^2 distribution. Hence one would not need to compute the second derivative of l in order to get Fisher information (or equivalently the asymptotic standard error of the MLE).

In order to solve the maximization problem, one often resolve to use the EM algorithm. ML still has the problem of undercoverage.

Assumptions Sample size has to be large enough for the ML estimates be approximately unbiased and normally distributed and with missing data, the sample might be larger than usual. Then likelihood functions comes from an assumed parametric model for complete data $P(Y_{obs}, Y_{mis}; \theta)$, hence departure from model assumptions might effect inference. MAR is still assumed.

3.2.2 Multiple imputation

Multiple imputation (MI) solves the problem of understating uncertainty. MI is similar to bootstrapping methods: one make artificial B samples and complete-case analysis. The final estimates (except standard errors) are then the arithmetic mean. Standard errors should reflect missing-data uncertainty and finite-sample variation.

An advantage of MI is the number of need imputation: the efficiency based on m samples relative o an infinite number is $(1 + \lambda/m)^{-1}$, where λ is the rate of missing information, which measures the increase in the large-sample variance of a parameter estimate due to missing values. $m = 20$ is often good in practice.

Combining standard errors In the one-dimensional case, if the sample is large enough so that the estimator Q follows a Gaussian distribution, then the estimate \hat{Q} and the standard error T can be computed from the estimates of $(Q^j, U^j)_{j=1}^m$, Q^j , respectively, U^j being the fitted value of Q , respectively the standard error, for data sets j

$$\begin{aligned}\hat{Q} &= m^{-1} \sum_{j=1}^m Q^j, \\ \hat{U} &= m^{-1} \sum_{j=1}^m U^j, \\ B &= (m-1)^{-1} \sum_{j=1}^m (Q^j - \hat{Q})^2, \\ T &= \hat{U} + (1 + m^{-1})B.\end{aligned}$$

For confidence interval, the Student's t approximation can be used with the degree of freedom given by

$$\nu = (m-1) \left[1 + \frac{\hat{U}}{(1 + m^{-1})B} \right]^2.$$

The estimated rate of missing information for Q is approximately $\tau/(\tau + 1)$ where $\tau = (1 + m^{-1})B/\hat{U}$, the relative increase in variance due to non-response. See [schafer1997@multivariate] for more cases.

This model still use the MAR assumption.

Obviously, the missing values problem is dealt before the analysis with MI, in contrast with ML. The danger from MI is the ability ot use different models for imputation and analysis.

4 Data analysis using regression and multilevel/hierarchical models [@gelman2006data]

Imputation of several missing data One could use start with multivariate regression with multivariate responses. The weakness of this method is the computational costs. However, one could use an iterative and cycling regression (like in GAM fitting) to assess the missing values.

A weakness of the iterative process is to make sure that all regression coefficient are consistent with each other.

5 Statistical analysis with missing data [little2002statistical]

The monograph describes mechanisms underlying the missingness come in several type (*mi*, *mice*, *Amelia* in R packages).

6 Imputing Missing Data for Gene Expression Array [hastie1999imputing]

Technical report reporting statistical methods for data imputation applied to human tumor data and subset of a subset yeast data [lichman2013yeast].

6.1 SVD

- Learn a set of basis functions (*eigen-genes*) from the complete-case analysis.
- Impute the missing cells for a gene by regressing its non-missing entries on the eigen-genes, and use the regression function to predict the expression values at the missing locations.
- The number of eigen-gene should be quite a bit smaller than the number of non-missing observations.

The data consists of $X \in \mathbb{R}^{N \times p}$, $N = 6830, p = 64$. X^c is the subset of complete genes (2069) and X^m the remainder.

6.1.1 SVD imputation using a clean training set

Singular value decomposition Remind that singular values of a matrix X are the square root of the non-negative eigenvalues of $X^T X$. Singular value decomposition (SVD) is provided by

$$\hat{X}_J^c = U_J D_J V_J^T, \quad (1)$$

where $D_J \in \mathbb{R}^{N \times p}$ is a diagonal matrix containing the leading $J < p$ singular values of X^c and $V_J \in \mathbb{R}^{p \times p}$ and $U_J \in \mathbb{R}^{N \times N}$, the corresponding orthogonal matrix of J right and left singular vectors. \hat{X}^c is the nearest matrix of X^c among matrices with rank J with respect to the sum of squares norm $\|A\|^2 = \text{tr}(AA^T)$.

Regression interpretation If x_i is any row of X^c , consider the regression of the p values in $x = (x_1, \dots, x_p)^T$ on the *eigen-gens* v_1, \dots, v_J , each p dimensional vectors. The regression solves

$$\min_{\beta} \|x - V_J \beta\|^2 = \min_{\beta} \sum_{l=1}^p \left(x_l - \sum_j 1^J v_{lj} \beta_j \right)^2 \quad (2)$$

with solution $\hat{\beta} = (V_J^T V_J)^{-1} V_J^T x = V_J^T x$ (since V_J is orthogonal) and orthogonal values $\hat{x}_l = V_l \hat{\beta}, l \in \{1, \dots, J\}$. Thus, according to Equation 1, $X^c V_J = U_J D_J$ gives all the

(transposed) regression coefficients for all the rows and $\hat{X}^c = U_J D_J V_J^T$ all the fitted values. Hence, once the matrix V_J is computed, SVD approximate each row of X^c by its fitted vector obtained by regression (or projection) on V_J . This suggest for a row of X^m with some missing components, they could possibly be imputed from

$$a \min_{\beta} \sum_{l=1}^p 1(M_l = 0) \left(x_l - \sum_j 1^J v_{lj} \beta_j \right)^2$$

where M_l is the missingness indicator of x_l

The imputation procedure is described as the following.

1. Compute the SVD of X^c and keep V_J .
2. For a row x^* with missing element, compute

$$\hat{\beta}^* = (V_J^{*T} V_J^*)^{-1} V_J^{*T} x^* a$$

where V_J^* is the *shortened* version of V_J with the appropriate rows removed (corresponding the missing elements of x^*). Note V_J^* no longer has orthogonal columns.

3. The predictions of the missing elements are $V_J^{(*)} \hat{\beta}^*$ where $V_J^{(*)}$ is the complement in V_J of V_J^* .

Usually, the data matrix is centered before SVD, however, for missing data, an intercept has to be fitted and a method based simulation is provided afterwards.

6.1.2 SVD imputation using all data

The previous methods usually discards a great number of data, particularly when $p \geq N$. In contrast, the next iterative procedure circumvent the problem at the cost of more computation.

1. Set X^* as X with all missing values filled by the mean of their row.
2. Solve the problem

$$\min_{V_J, D_J, U_J} \|X^* - m 1^T - U_J D_J V_J^T\|^* \quad (3)$$

where $\|\cdot\|^*$ is the sum of squares of all non-missing elements and $m \in \mathbb{R}^N$ is the row means of X^* .

3. Predict the missing values of X with the fitted values.
4. Reset X^* as X with the missing values replaced by the result of previous step.
5. Repeat step 2-5, until the size of the relative update of the missing values become negligible.

The paper states that only 6 iterations is necessary. An interesting point is the solution of Equation (3) is a fixed point, i.e. if missing values are filled, and the SVD algorithm is executed on the “complete” matrix, the solution remain the same. Additionally, if the *eigen-genes* obtained from Equation (3), and impute the missing value using the regression approach in (2), the imputations are identical to those obtained from (3).

6.2 K-nearest neighbor averaging imputation

The paper described the other end of the spectrum in term of data usage *K-nearest neighbor averaging*. The algorithm is described as following.

1. Computed the Euclidean distance between x^* and all the genes in X^c , using only those co-ordinates not missing in x^* . Identify the K closest.
2. Impute the missing coordinates of x^* by averaging the corresponding coordinates of the K closest.

The authors then found that K between 5 to 10 was a good choice for their data set.

6.3 Repeated regression

This algorithm iteratively impute the missing values of each column, by making a regression against all the other columns (with imputed values). The method iterates until convergence of the filled values.

The number of iteration to convergence is typically materialy, constituting a legible drawback.

It differentiates itself to the SVD by making the regression directly on the data matrix. The author claim that CARTs can replaced the regression and avoid avoid iteration.

7 Missing value estimation methods for DNA microarrays [atroyanskaya2001missing]

This should be a more elaborated papers than the previous paper.

1. SVD based method (SVNimpute)
2. Weighted K-nearest neighbors (KNNimpute).
3. Row average.

KNNimpute provides supress SVDimpute in simulation with missing rate between 1 – 20%.

In microarray studies, rows are genes expression and columns are different experimental conditions. Missing data are common, and there are no obvious causes for this missingness. Row average is does not take correlation into account.

Non-response issues in sample surveys and missing data in experiments [alittle2002statistical].

Three types of data are studied: noisy time series, time series and non-time series. Each data sets is cleaned to yield *complete* matrices.

A missing rate p is set between 0.01 and 0.2 and deleted at random for each data sets, from which each recovering methods is used to fill the missing values. The fitted values are then compared to the original values using the normalized root mean squared (normalized RMS).

KNNImpute The imputed value is provided by weighted value of the K most similar neighbor. The Euclidean distance measure was the often the most accurate of measure similarity. The normalized weights are proportional to the inverse of the distance. Transforming the data provides additional robustness.

At least four columns are necessary to have a decent performance. Complexity is $O(N^2p)$

SVDImpute The method is the same as the one described in [hastie1999imputing]. Complexity is $O(Np^2i)$, where i is the number of iterations for the convergence of the SVD algorithm (5 to 6).

Conclusion KNN is better suited than SVD for non-time series or noisy data and is more robust to the increasing missing rates.

8 Datasets

- Use the data sets in the appendix of [schafer1997analysis], FLAS inside the `miP` package.
- Use the synthetic bivariate set of [schafer2002missing], (normal and cluster).
- Data sets: titanic, iris, mtcars, yeast [lichman2013yeast], (noisy) time series data?

9 Chained Equations

The goal is to use Gibbs sampling to generate observations by making p sampling, using $Y_j|Y_{-j}, \theta$.

Predictive mean matching Statistical horizon explains how predictive mean matching (pmm) works. It is similar to nearest neighbor in the sense that predictions from linear regression is used as the metric and the imputed value is randomly chosen by the sets of neighbor. This is quite practical as the imputed observations always belong to the set of non-missing data.

10 FLAS Completion

The `mi` data package provides imputed continuous data, with decimals (e.g. 102.3435), whereas the set of observed data is clearly the integers.

10.1 MI

MI has a problem of convergence when using correlated data example with `grades` and `grades_complete`.

11 Bibliography